

A decorative graphic on the left side of the slide consisting of two overlapping parallelograms. The front one is blue and the back one is a light green color. They are positioned diagonally, with the blue one partially covering the green one.

# What is Kubernetes?

## An intro to container orchestration

With Stephen Tollenaar



# A background

- DevOps Engineer at CarteNav
- 6 years of Kubernetes experience
- Certified Kubernetes Security Specialist
- Small homelab



# Before we dive in

- No expectations of participation
  - But if you want use one of these on your laptop
    - Minikube
    - Microk8s
    - DO NOT USE DOCKER SWARM
- Feel free to ask questions

# Quick reminder

```
FROM node:22 (last pushed 6 hours ago)
WORKDIR /app
COPY package*.json ./
RUN npm install
COPY . .
EXPOSE 5000
CMD ["node", "index.js"]
```

So we know what Docker is

```
stollenaar@spicetop:~/../personal/kubetalk$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
f56b4d1338c1	fullstack-frontend	"docker-entrypoint.s..."	3 seconds ago	Up 3 seconds	0.0.0.0:3000->80/tcp, [::]:3000->80/tcp	fullstack-frontend-1
7a80e55ceb46	fullstack-backend	"docker-entrypoint.s..."	3 seconds ago	Up 3 seconds	0.0.0.0:5000->5000/tcp, [::]:5000->5000/tcp	fullstack-backend-1
4cdf50862ebd	postgres:15	"docker-entrypoint.s..."	2 minutes ago	Up 3 seconds	5432/tcp	fullstack-db-1
4760aef03187	moby/buildkit:buildx-stable-1	"buildkitd"	9 months ago	Up 4 days		buildx_buildkit_competent_wescoff0

```
stollenaar@spicetop:~/../personal/kubetalk$
```

# Full Stack App

```
version: "3.9"

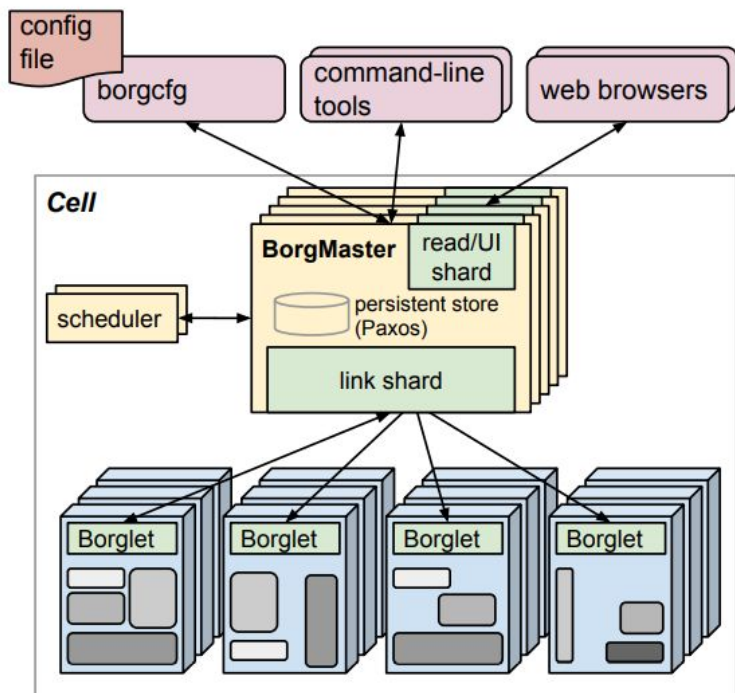
services:
  frontend:
    build: ./frontend
    ports:
      - "3000:80"
    depends_on:
      - backend

  backend:
    build: ./backend
    ports:
      - "5000:5000"
    environment:
      - DATABASE_URL=postgres://postgres:postgres@db:5432/postgres
    depends_on:
      - db

  db:
    image: postgres:15
    restart: always
    environment:
      POSTGRES_USER: postgres
      POSTGRES_PASSWORD: postgres
      POSTGRES_DB: postgres
    volumes:
      - db_data:/var/lib/postgresql/data

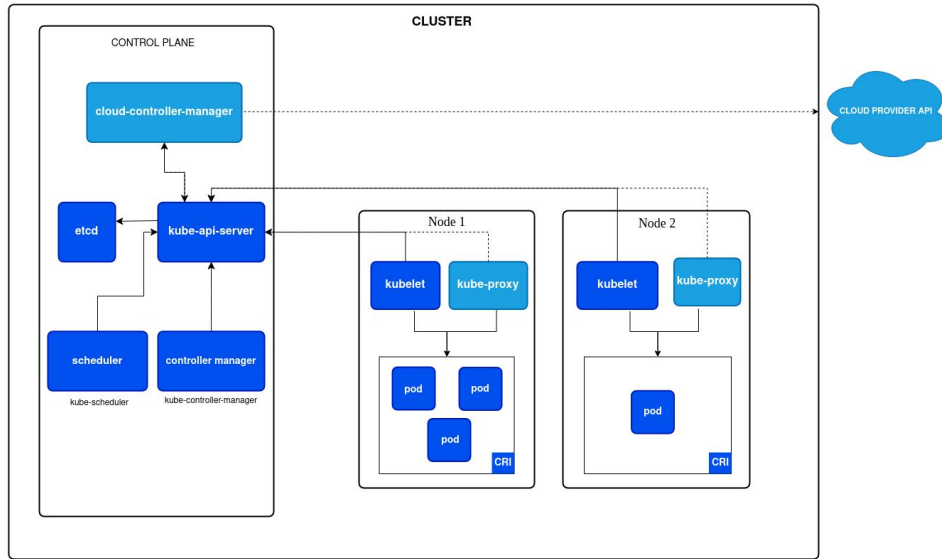
volumes:
  db_data:
```

# Borg



- The first iteration of container orchestration
- Does feature “pods”
- Borg tasks are run inside a Linux cgroup-based resource container
- Still very open to accessing host resources

# Introducing Kubernetes



- Removes the headache of managing containers
  - Your sanity might falter
- You specify what you want, Kubernetes takes care of the rest
- Scaling up to an insane number of resources\*
  - No more than 110 pods per node
  - No more than 5,000 nodes
  - No more than 150,000 total pods
  - No more than 300,000 total containers

# What is a Pod?

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.29.1
      ports:
        - containerPort: 8080
```

- The lowest of building blocks
- The container array is “similar” to the docker compose services
- Can hold more containers, and even init containers
- We apply/delete these manifests using kubectl

```
version: "3.9"
services:
  nginx:
    image: nginx:1.29.1
    ports:
      - "8080:8080"
```





# Pod Deployment methods

But a pod is just a singular throwaway thing, we need something more “persistent”

So we got:

- Deployment
  - ReplicaSet
- StatefulSet
- DaemonSet
- Job
- CronJob

# What is a Deployment?

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.29.1
          ports:
            - containerPort: 8080
```

- A Deployment manages a set of Pods to run an application workload, usually one that doesn't maintain state.
- The template resembles the pod from before
- Each version of a Deployment creates a ReplicaSet

```
stollenaar@spicetop:~/.../personal/kubetalk$ kubectl get replicaset.apps
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-844c99f596	3	3	0	2s

```
stollenaar@spicetop:~/.../personal/kubetalk$
```

```
stollenaar@spicetop:~/.../personal/kubetalk$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-844c99f596-cdbb4	1/1	Running	0	76s
nginx-844c99f596-lxq5d	1/1	Running	0	76s
nginx-844c99f596-rm9l6	1/1	Running	0	76s

```
stollenaar@spicetop:~/.../personal/kubetalk$
```

# What is a StatefulSet?

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: postgres
spec:
  serviceName: "postgres"
  replicas: 1
  selector:
    matchLabels:
      app: postgres
  template:
    metadata:
      labels:
        app: postgres
    spec:
      containers:
        - name: postgres
          image: postgres:17
          ports:
            - containerPort: 5432
          env:
            - name: POSTGRES_PASSWORD
              value: password
          volumeMounts:
            - name: pgdata
              mountPath: /var/lib/postgresql/
      volumeClaimTemplates:
        - metadata:
            name: pgdata
          spec:
            accessModes: ["ReadWriteOnce"]
            resources:
              requests:
                storage: 1Gi
```

- A StatefulSet runs a group of Pods, and maintains a sticky identity for each of those Pods.
- Useful for databases.
- Same building block from the Pod
- No ReplicaSet
- Has the replica number, instead of hash

```
stollenaar@spicetop:~/.../personal/kubetalk$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
postgres-0    1/1     Running   0           89s
stollenaar@spicetop:~/.../personal/kubetalk$
```

# What is a DaemonSet?

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: node-exporter
  labels:
    app: node-exporter
spec:
  selector:
    matchLabels:
      app: node-exporter
  template:
    metadata:
      labels:
        app: node-exporter
    spec:
      containers:
        - name: node-exporter
          image: prom/node-exporter:latest
          ports:
            - containerPort: 9100
```

- A DaemonSet ensures that all (or some) Nodes run a copy of a Pod

Some typical uses of a DaemonSet are:

- Running a cluster storage daemon on every node
- Running a logs collection daemon on every node
- Running a node monitoring daemon on every node

# What is a Job?

```
apiVersion: batch/v1
kind: Job
metadata:
  name: hello-job
spec:
  template:
    spec:
      containers:
      - name: hello
        image: busybox
        command: ["echo", "Hello, Kubernetes!"]
      restartPolicy: Never
```

- Jobs represent one-off tasks that run to completion and then stop.

# What is a CronJob?

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: postgres-backup
spec:
  schedule: "0 2 * * *" # Runs daily at 2 AM
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: pg-backup
              image: postgres:16
              env:
                - name: PGPASSWORD
                  value: examplepassword
              command:
                - /bin/sh
                - -c
                - pg_dump -h postgres -U postgres -F c -b -v -f /backup/backup.dump postgres
              volumeMounts:
                - name: backup-volume
                  mountPath: /backup
          restartPolicy: OnFailure
```

- A regular recurring Job
- Builds on top of the Job template from before

# Where in Kubernetes do we run these things?

```
apiVersion: v1
kind: Namespace
▼ metadata:
  name: kubetalk
```

- Used for isolating groups of resources within a single cluster
- Not everything can be namespaced

```
stollenaar@spicetop:~/.../personal/kubetalk$ kubectl get ns
```

NAME	STATUS	AGE
cert-manager	Active	406d
cmstate-operator	Active	406d
copypastabotv2	Active	258d
default	Active	406d
diplomacy	Active	214d
external-dns	Active	23d
github-arc	Active	128d
jellyfin	Active	406d
kube-node-lease	Active	406d
kube-public	Active	406d
kube-system	Active	406d
kubelet-serving-cert-approver	Active	367d
kubetalk	Active	4d22h
metallb-system	Active	406d
monitoring	Active	361d
ollama	Active	227d
ollamabot	Active	41d
openebs	Active	406d
renovate	Active	102d
statisticsbot	Active	406d
tailscale	Active	406d
uptime-kuma	Active	205d
vault	Active	406d

```
stollenaar@spicetop:~/.../personal/kubetalk$
```



# The Kubernetes Node

- Runs your work loads
- Is either a physical or virtual workload

```
stollenaar@spicetop:~/.../personal/kubetalk$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
talos-7zr-i5q	Ready	control-plane	197d	v1.34.0
talos-e5t-zk5	Ready	worker	406d	v1.34.0
talos-iso-cgi	Ready	worker	266d	v1.34.0

```
stollenaar@spicetop:~/.../personal/kubetalk$
```



# Putting it all Together

```
stollenaar@spicetop:~$ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/kubetalk-app-6bb5d66f99-5scfb	1/1	Running	0	2m34s
pod/kubetalk-backend-664c48448c-zwg5w	1/1	Running	0	2m34s
pod/postgres-0	1/1	Running	0	3m13s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubetalk-app	ClusterIP	10.104.95.93	<none>	80/TCP	3m13s
service/kubetalk-backend	ClusterIP	10.106.152.120	<none>	4000/TCP	3m13s
service/postgres	ClusterIP	10.106.230.46	<none>	5432/TCP	3m13s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/kubetalk-app	1/1	1	1	2m34s
deployment.apps/kubetalk-backend	1/1	1	1	2m34s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/kubetalk-app-6bb5d66f99	1	1	1	2m34s
replicaset.apps/kubetalk-backend-664c48448c	1	1	1	2m34s

NAME	READY	AGE
statefulset.apps/postgres	1/1	3m13s

```
stollenaar@spicetop:~$
```

- The fullstack app in a namespace

# Another Example

```
stollenaar@spicetop:~/.../personal/kubetalk$ kubectl get pods,deployments,statefulset
```

NAME	READY	STATUS	RESTARTS	AGE
pod/bazarr-768c7d8d78-lthf6	1/1	Running	1 (3d8h ago)	4d
pod/byparr-6c9c5b8749-tlxts	1/1	Running	3 (10h ago)	5d8h
pod/jellyfin-5847f9896d-8b52h	1/1	Running	0	47h
pod/jellyseerr-b644f77d-8gp4r	1/1	Running	0	3d13h
pod/postgres-1	1/1	Running	0	3d3h
pod/prowlarr-5fffb576bc-2jw47	1/1	Running	1 (3d8h ago)	4d
pod/qbittorrent-5c77b5d657-5q2lj	2/2	Running	0	3d3h
pod/radarr-d7c9bc8d5-v674v	1/1	Running	1 (3d8h ago)	4d
pod/sonarr-df775568c-fkkbz	1/1	Running	1 (3d8h ago)	4d

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/bazarr	1/1	1	1	121d
deployment.apps/byparr	1/1	1	1	75d
deployment.apps/decluttarr	0/0	0	0	127d
deployment.apps/jellyfin	1/1	1	1	416d
deployment.apps/jellyseerr	1/1	1	1	416d
deployment.apps/prowlarr	1/1	1	1	416d
deployment.apps/qbittorrent	1/1	1	1	253d
deployment.apps/radarr	1/1	1	1	416d
deployment.apps/sonarr	1/1	1	1	416d
deployment.apps/tdarr-node	0/0	0	0	16d
deployment.apps/tdarr-server	0/0	0	0	16d



But what about extra config?

# What is a ConfigMap?

```
vi@configmap+vi@configmap.json
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-env
data:
  DB_HOST: localhost
  DB_PORT: "5432"
  NODE_ENV: production
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-file
data:
  user-interface.properties: |
    color.good=purple
    color.bad=yellow
    allow.textmode=true
```

- A ConfigMap is used to store non-confidential data in key-value pairs.
- Pods can consume ConfigMaps as environment variables, command-line arguments, or as configuration files in a volume.
- When values are updated, it automatically propagates

# What is a Secret?

```
apiVersion: v1
kind: Secret
metadata:
  name: app-secret
  namespace: kubetalk
type: Opaque
stringData:
  POSTGRES_PASSWORD: examplepassword
```

- A Secret is an object that contains a small amount of sensitive data such as a password, a token, or a key
- Are similar in config as a ConfigMap
- When viewing the values, they are encoded

```
stollenaar@spicetop:~$ kubectl get secret app-secret -o yaml
apiVersion: v1
data:
  POSTGRES_PASSWORD: ZXhhbXBsZXBhc3N3b3Jk
kind: Secret
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","kind":"Secret","metadata":{"annotations":{},"name":"app-secret","namespace":"kubetalk"},"stringData":{"POSTGRES_PASSWORD":"examplepassword"}}
  creationTimestamp: "2025-09-16T23:28:34Z"
  name: app-secret
  namespace: kubetalk
  resourceVersion: "89063781"
  uid: bd1fae01-b19c-4c57-8cea-d720ff2331dd
type: Opaque
```



# But what about storage?

- We have persistent storage and ephemeral

# What is a PersistentVolume?

```
io.k8s.api.core.v1.PersistentVolume (v1@persistentvolume.json)
apiVersion: v1
kind: PersistentVolume
metadata:
  annotations:
    pv.kubernetes.io/provisioned-by: nfs.csi.k8s.io
    volume.kubernetes.io/provisioner-deletion-secret-name: ""
    volume.kubernetes.io/provisioner-deletion-secret-namespace: ""
  finalizers:
    - kubernetes.io/pv-protection
  name: pvc-d7c4c439-916a-4476-88c2-24a84429a485
spec:
  accessModes:
    - ReadWriteOnce
  capacity:
    storage: 20Gi
  claimRef:
    apiVersion: v1
    kind: PersistentVolumeClaim
    name: postgres-1
    namespace: jellyfin
    resourceVersion: "89451732"
  csi:
    driver: nfs.csi.k8s.io
    volumeAttributes:
      csi.storage.k8s.io/pv/name: pvc-d7c4c439-916a-4476-88c2-24a84429a485
      csi.storage.k8s.io/pvc/name: postgres-1
      csi.storage.k8s.io/pvc/namespace: jellyfin
      server: 192.168.2.113
      share: /mnt/main/kubernetes
      storage.kubernetes.io/csiProvisionerIdentity: 1758195604893-1373-nfs.csi.k8s.io
      subDir: jellyfin/postgres-1
    volumeHandle: 192.168.2.113#mnt/main/kubernetes#jellyfin/postgres-1#pvc-d7c4c439-916a-4476-88c2-24a84429a485#
  mountOptions:
    - noLOCK
    - vers=4
  persistentVolumeReclaimPolicy: Retain
```

- Describes in detail what your storage for your volume is

# What is a PersistentVolumeClaim?

io.k8s.api.core.v1.PersistentVolumeClaim (v1@persistentvolumeclaim.json)

apiVersion: v1

kind: PersistentVolumeClaim

metadata:

annotations:

cnpg.io/nodeSerial: "1"

cnpg.io/operatorVersion: 1.27.0

cnpg.io/pvcStatus: ready

pv.kubernetes.io/bind-completed: "yes"

pv.kubernetes.io/bound-by-controller: "yes"

volume.beta.kubernetes.io/storage-provisioner: nfs.csi.k8s.io

volume.kubernetes.io/storage-provisioner: nfs.csi.k8s.io

creationTimestamp: "2025-09-18T19:50:43Z"

finalizers:

- kubernetes.io/pvc-protection

labels:

cnpg.io/cluster: postgres

cnpg.io/instanceName: postgres-1

cnpg.io/instanceRole: primary

cnpg.io/pvcRole: PG\_DATA

role: primary

name: postgres-1

namespace: jellyfin

ownerReferences:

- apiVersion: postgresql.cnpg.io/v1

controller: true

kind: Cluster

name: postgres

uid: 90afed56-3ble-4c9c-92df-1cc4252ad04b

spec:

accessModes:

- ReadWriteOnce

resources:

requests:

storage: 20Gi

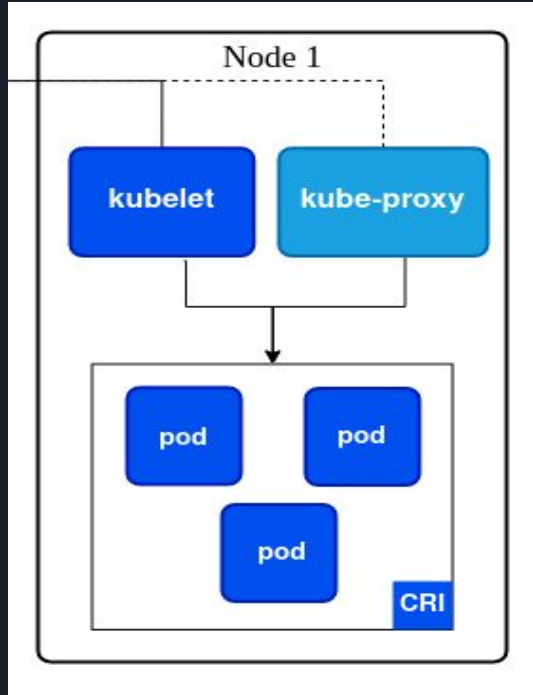
storageClassName: nfs-csi-main

volumeMode: Filesystem

- Tells kubernetes what you want, how large, and from which provider
- Interacts with the StorageClass, which handles the volume requests



# But How do they talk to each other?



- The easy: magic
- The complicated: Container Network Interfaces
- The popular:
  - Flannel
  - Calico
  - Cilium
- The cloud platforms usually use their own
- Runs on each node, (Daemonset)

# The Services

```
io.k8s.api.core.v1.Service (v1@service.json)
apiVersion: v1
kind: Service
metadata:
  name: sonarr
  namespace: jellyfin
spec:
  ports:
    - port: 8989
      protocol: TCP
      targetPort: 8989
  selector:
    app: sonarr
  sessionAffinity: None
  type: ClusterIP
```

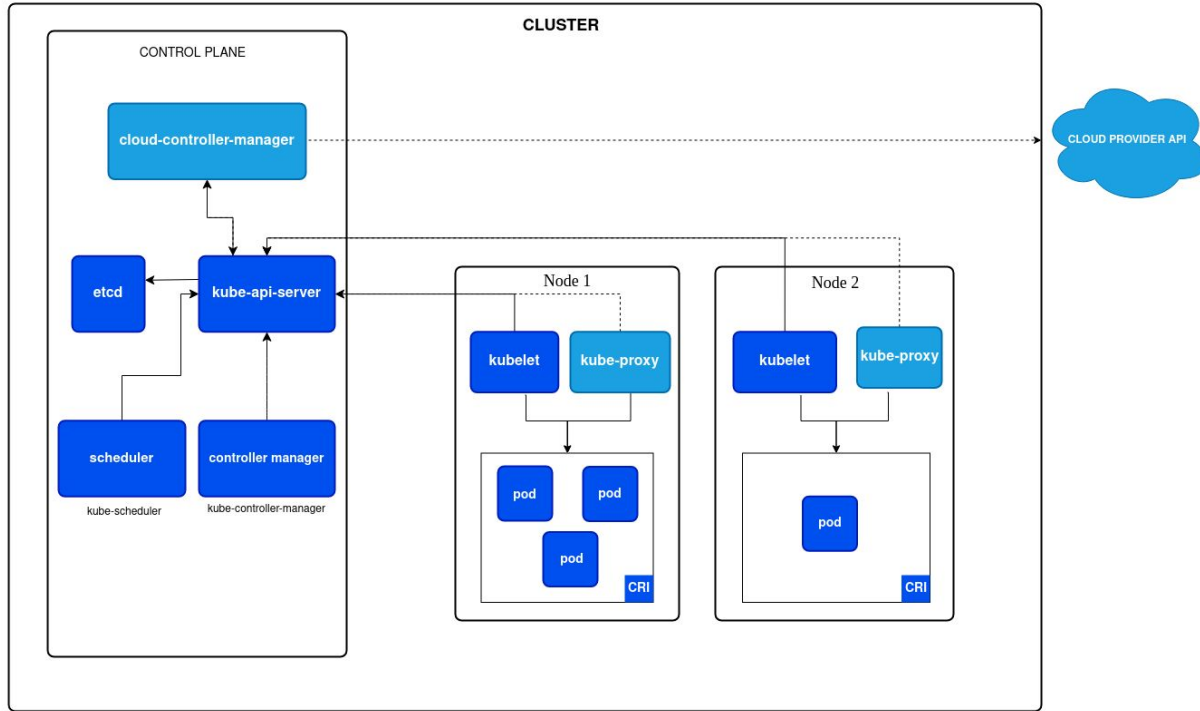
- Since Pods come and go, so do their Pod IPs
- We need something that stays the same so other Pods don't crash
- Knows where to send the network traffic through the labels
- 4 types:
  - ClusterIP
    - For internal traffic
  - LoadBalancer
    - To external incoming traffic in a balancing method
  - NodePort
    - To handle incoming traffic if loadbalancing is not an option
  - ExternalName
    - CNAME for a DNS

# The Ingress

```
io.k8s.api.networking.v1.Ingress (v1@ingress.json)
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  annotations:
    cert-manager.io/cluster-issuer: letsencrypt-prod
    kubernetes.io/ingress.class: nginx
  name: jellyfin
  namespace: jellyfin
spec:
  ingressClassName: nginx
  rules:
    - host: jellyfin.home.spicedelver.me
      http:
        paths:
          - backend:
              service:
                name: jellyfin-web
                port:
                  number: 8096
              path: /
              pathType: ImplementationSpecific
      tls:
        - hosts:
            - jellyfin.home.spicedelver.me
          secretName: jellyfin-tls
```

- Make your HTTP (or HTTPS) network service available using a protocol-aware configuration mechanism, that understands web concepts like URIs, hostnames, paths, and more.
- You need a controller to direct this kind of traffic

# But what is Kubernetes Actually?



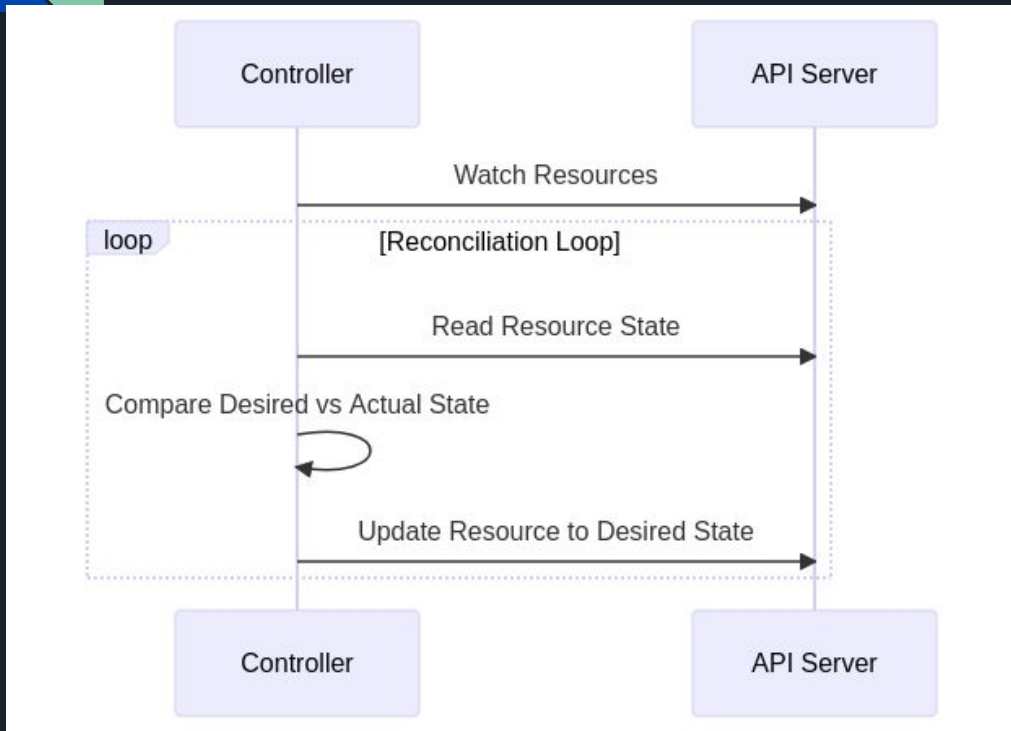
- Everything is a controller, It's controllers all the way down
- Reconciliation loops



# The Controllers

- Kube-api-server
  - REST API to manage EVERYTHING in a cluster
- Kubelet
  - Manages the node within a cluster
- Scheduler
  - Makes sure pods can be provisioned and run
- Controller-manager
  - Makes sure the desired state of a cluster is achieved
- ETCD/DQLite
  - The Key Value store of all your cluster resources
- Kube-Proxy
  - Manages your container network, this will differ per cluster

# How does a controller work



- The Reconciliation loop
- We introduce CustomResourceDefinitions to expand what a Resource is



# Custom Controller Example

- Quick demo from my cmstate-operator  
<https://github.com/STollenaar/cmstate-injector-operator/tree/main>



# PizzaKube

Yes, kubernetes could order you a pizza:

<https://github.com/grantgumina/provider-pizza/blob/master/examples/order/example-order.yaml>





# Questions?



# Resources

- Find me almost every Thursday evening at CTS
  - 7pm-9pm Jumping Bean Elizabeth Ave
- <https://github.com/STollenaar/kubetalk>

